

ZARGAR · TIPS TECHNIQUE

The Tips Knowledge System


How the desk reads its Discord rooms, remembers what matters, forgets what doesn't — and how we test the whole loop against history before trusting it with money.



Built 2026-08-30 · docs/techniques/tip/KNOWLEDGE-PLAN.md (research & decisions) · KNOWLEDGE-BUILD-PLAN.md (build & findings)

Research foundation: Zep temporal knowledge graphs · FinMem layered memory · FinCon per-agent isolation (deep-research run, 23/25 claims verified)

1 · The big picture — from chatter to knowledge



Two analysts, two brains: the Tips Analyst's knowledge lives in tip_notes; EM's lives in its own rulebook/prompts. A static guard test fails the suite on any cross-read (PLATFORM-RULES invariant 12).

2 · Knowledge that ages honestly

The research is blunt: stale memory misleads models, and similarity search can't tell a contradiction from a rephrase. So knowledge here has a **lifecycle** — layered retention (FinMem), temporal invalidation instead of deletion (Zep), LLM-judged audits.


Scope	What lives there	Retention	How it dies
rule	the analyst's own trading rules (numbered, injected into every run)	NO EXPIRY	weekly audit merges/expires; contradictions need your click
ticker:SYM	structural reads on one name	90 DAYS	expires unless cited in a live run (refresh) or 📌 pinned
source:name	a tipper's habits & tells	90 DAYS	same — citation refreshes, audit consolidates
general	desk-wide context	NO EXPIRY	audit-gated
daily:date	digest of a context channel's day	14 DAYS	just expires — chatter is short-lived by design
experiment:b	batch findings	NEVER INJECTED	review artifacts, out-of-band forever

**Citation refresh**
A note injected into a **completed live appraisal** gets `cited_count++` and its TTL extended. Knowledge that keeps participating in decisions stays alive; knowledge nobody uses ages out. Historical experiments deliberately never refresh.

**Query-time expiry**
Nothing is deleted or swept. An expired note simply stops being served — deterministic, restart-proof, and fully visible under the Knowledge tab's history toggle (superseded + expired).

**Widened audit**
Weekly, the same judge→apply pass runs over the rulebook *and* every ticker/source/general group with ≥3 notes: merge near-duplicates, expire the unsupported, flag real contradictions **needs-your-call**.


Digests: the trading-floor pipeline



Digest-now button per context channel (Tips → Sources) tunes the prompt first; the nightly pass turns on only when `techniques.tip.digest_enabled` flips.

3 · The historical experiment — testing the desk against the past

The published field has no rigorous answer for replay-testing an LLM tips desk — so this loop *is* the evidence. It grades the **process** (extraction, verification, knowledge, tools, reasoning), never the P&L: hindsight is cheap, process is fixable.



Running a batch

```
python -m zargar.tools.tip_experiment run --batch b1 --sample 20 --seed 7 --since 2026-06-01
python -m zargar.tools.tip_experiment status --batch b1
python -m zargar.tools.tip_experiment review --batch b1
```

Batch 1 — ran 2026-08-30

20 messages → 22 signals, all replayed · **0 orders, 0 proposals** (verified live) · 19 appraised, all skip · **12 findings** (3 silent drops, freshness/ticker checks acting as constants, premium-vs-underlying confusion, post-tip leakage...) — the harness found more bugs in one afternoon than weeks of live trickle. Full list: KNOWLEDGE-BUILD-PLAN.md Phase 5.

Batch 2 → recurrence check

After the F1-F12 fixes: a fresh seed, same rubric — did the same failure classes come back? Resolved classes get ticked, recurring ones escalate.

Then monthly

A batch a month keeps the pipeline honest as sources, prompts and rules evolve.

4 · The nightly routine — what to check each evening

When (ET)	What runs	What you look at
16:45	Flow scan (chain snapshots → reads)	Flow page: new reads, universe joins
evening	Nightly tip review: retros on closed positions, unfilled retros, lane grades, digests (when enabled)	Tips → Analyst: retro runs; Knowledge → 📅 Today: is the digest a fair read of the room?
audit day	Rule audit + knowledge audit (weekly)	Knowledge → ⚠ Needs you: resolve contradictions — keep one, delete one, or keep both deliberately
any	—	Approvals: nothing pending unseen · Armed: plans healthy, no needsAttention

Your 5-minute knowledge pass

- 📅 **Read today's digest(s).** Wrong lean or missed theme → tune with digest-now before enabling nightly.
- ⚠ **Clear needs-your-call flags.** They only appear when two notes genuinely pull opposite ways.
- ⌚ **Glance at expiring notes.** Amber chip (≤7d): still true? 📌 pin it. Stale? let it die — that's the system working.
- 🧪 **Experiment cadence.** After fixes land, queue the next batch with a new seed; findings go to the build plan.
- 📌 **Teach one thing.** Anything you learned watching the rooms today belongs in the composer — scoped, so only the runs that need it see it.

Healthy signs

Rules stay few and sharp · digests promote 0-2 nuggets a day, not 5 · expiring notes mostly deserve to expire · batch reviews find *new* classes, not repeats.

Smells

Rulebook growing weekly · every note pinned · promotions restating the daily summary · the same extraction bug in two batches — stop and fix before more volume.

5 · How we got here — the thought process

2026-08-27 · The platform bet

Many techniques, one engine. The Tip technique becomes the second resident: extraction → verification → dual shadow books per source, so every tipper earns (or loses) trust on their own record.

2026-08-28 · The analyst is born

Tips deserve judgement, not just parsing: an independent trader persona with market tools, its own exit plans, its own rulebook — explicitly NOT bound by EM's method book.

2026-08-29 · Staying armed, closing gaps

Multi-day stay-armed plans, follow-up handling ("a close never opens"), never-chase caps, nightly retros + the first weekly rule audit. The desk starts maintaining its own judgement.

2026-08-30 · Knowledge grows a lifecycle

Deep research (Zep, FinMem, FinCon + 20 more sources, adversarially verified) settles the design: typed scoped stores, temporal expiry over deletion, citation-refresh, LLM-judged audits, per-agent isolation. Context channels + digests bring trading-floor in without polluting intake. The experiment harness closes the loop: history in, findings out, fixes in, repeat.

2026-08-30 · Batch 1 proves the loop

First historical batch (20 messages, seed 7): the money guarantees held perfectly (0 orders, 0 books), the judgment layer graded well ("that is real process"), and the rubric review surfaced 12 concrete findings — silent drops, checks acting as constants, unit confusion, post-tip leakage. Exactly what the experiment was built to catch, caught on day one.

Next

F1-F12 fixes → batch 2 with a fresh seed · nightly digests once tuned · monthly experiment cadence · knowledge volume grows → revisit retrieval (the temporal-graph option stays on the shelf until scale demands it).

Design principles (the short version)



Process over hindsight

Grade what the desk could have known, never what the chart did later.



Scope everything

Knowledge reaches only the runs it belongs to; two analysts never share a brain.



Forget on purpose

Expiry + audits keep the context honest; only cited or pinned knowledge earns permanence.



Experiments can't spend

Out-of-band by construction, proven by tests — curiosity never risks a dollar.